

Classifying CIFAR-10 dataset with the help of Convolutional Neural Networks (CNN):

Import required libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

%matplotlib inline
```

Import CIFAR-10 dataset

```
In [2]: from tensorflow.keras.datasets import cifar10
```

```
In [3]: (X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170500096/170498071 [=====] - 32s 0us/step

Checking shape of training and testing data


```
In [4]: print('Shape of Dataset: ')
print(X_train.shape, '\t', X_test.shape, '\t', y_train.shape, '\t', y_test.shape)
```

Shape of Dataset:
(50000, 32, 32, 3) (10000, 32, 32, 3) (50000, 1) (10000, 1)

Display some sample images of the dataset


```
In [25]: from PIL import Image
Image.open('cifar-10.png')
```

Out[25]:




```
In [8]: plt.imshow(X_train[0])
```

```
Out[8]: <matplotlib.image.AxesImage at 0x24d4d2bb148>
```




```
In [9]: plt.imshow(X_train[4])
```

```
Out[9]: <matplotlib.image.AxesImage at 0x24d4d36d248>
```




```
In [10]: plt.imshow(X_train[7])
```

```
Out[10]: <matplotlib.image.AxesImage at 0x24d4d3d5048>
```



```
In [12]: plt.imshow(X_train[9])
```

```
Out[12]: <matplotlib.image.AxesImage at 0x24d4d49b5c8>
```



Checking max and min size and scaling the dataset

```
In [13]: X_train.max()
```

Out[13]: 255

In [14]: `X_train.min()`

Out[14]: 0

In [15]: `X_train = X_train/255`
`X_test = X_test/255`

In [19]: `y_train[0:11]`

Out[19]: array([[6],
 [9],
 [9],
 [4],
 [1],
 [1],
 [2],
 [7],
 [8],
 [3],
 [4]], dtype=uint8)

Convert y to categorical values: One hot encoding

In [20]: `from tensorflow.keras.utils import to_categorical`

In [21]: `y_cat_train = to_categorical(y_train, 10)`
`y_cat_test = to_categorical(y_test, 10)`

In [22]: `y_cat_train[0:11]`

Out[22]: array([[0., 0., 0., 0., 0., 0., 1., 0., 0., 0.],
 [0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
 [0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
 [0., 0., 0., 0., 1., 0., 0., 0., 0., 0.],
 [0., 1., 0., 0., 0., 0., 0., 0., 0., 0.],
 [0., 1., 0., 0., 0., 0., 0., 0., 0., 0.],
 [0., 0., 1., 0., 0., 0., 0., 0., 0., 0.],
 [0., 0., 0., 0., 0., 0., 0., 1., 0., 0.],
 [0., 0., 0., 0., 0., 0., 0., 0., 1., 0.],
 [0., 0., 0., 1., 0., 0., 0., 0., 0., 0.],
 [0., 0., 0., 0., 1., 0., 0., 0., 0., 0.]], dtype=float32)

Import libraries for building the CNN

In [26]: `from tensorflow.keras.models import Sequential`
`from tensorflow.keras.layers import MaxPool2D, Conv2D, Flatten, Dense`

Build the model

In [29]: `model = Sequential()`

`#Add Convolution Layer`
`model.add(Conv2D(filters=32, kernel_size=(4,4), input_shape = (32,32,3), activation='re`
`#Add Pooling Layer`
`model.add(MaxPool2D(pool_size=(2,2)))`

```

#Pixels/Total no of values inside a single MNIST image = 28*28 = 784
#Pixels/Total no of values inside a single CIFAR-10 image = 32*32*3 = 3072
#Since there is a lot more information present, it is better to add more Conv Layers to

#Add Convolution Layer
model.add(Conv2D(filters=32, kernel_size=(4,4), input_shape = (32,32,3), activation='re
#Add Pooling Layer
model.add(MaxPool2D(pool_size=(2,2)))

model.add(Flatten())

model.add(Dense(256,activation='relu'))

model.add(Dense(10,activation='softmax'))

model.compile(loss='categorical_crossentropy',optimizer='adam', metrics=['accuracy'])

```

View model summary

In [30]: `model.summary()`

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 29, 29, 32)	1568
max_pooling2d_4 (MaxPooling2)	(None, 14, 14, 32)	0
conv2d_5 (Conv2D)	(None, 11, 11, 32)	16416
max_pooling2d_5 (MaxPooling2)	(None, 5, 5, 32)	0
flatten_2 (Flatten)	(None, 800)	0
dense_4 (Dense)	(None, 256)	205056
dense_5 (Dense)	(None, 10)	2570
=====		
Total params: 225,610		
Trainable params: 225,610		
Non-trainable params: 0		

Adding Early Stopping

In [31]: `from tensorflow.keras.callbacks import EarlyStopping`

In [32]: `early_stop = EarlyStopping(monitor='val_loss', patience=2)`

Train the model

In [33]: `model.fit(X_train, y_cat_train, epochs=20, callbacks = [early_stop], validation_data=(X`

```

Train on 50000 samples, validate on 10000 samples
Epoch 1/20
50000/50000 [=====] - 82s 2ms/sample - loss: 1.5065 - accuracy:
0.4553 - val_loss: 1.2682 - val_accuracy: 0.5467
Epoch 2/20
50000/50000 [=====] - 61s 1ms/sample - loss: 1.1653 - accuracy:
0.5867 - val_loss: 1.0939 - val_accuracy: 0.6142

```

```

Epoch 3/20
50000/50000 [=====] - 60s 1ms/sample - loss: 1.0066 - accuracy:
0.6474 - val_loss: 1.0839 - val_accuracy: 0.6196
Epoch 4/20
50000/50000 [=====] - 59s 1ms/sample - loss: 0.8986 - accuracy:
0.6865 - val_loss: 1.0611 - val_accuracy: 0.6235
Epoch 5/20
50000/50000 [=====] - 57s 1ms/sample - loss: 0.8018 - accuracy:
0.7186 - val_loss: 0.9956 - val_accuracy: 0.6623
Epoch 6/20
50000/50000 [=====] - 57s 1ms/sample - loss: 0.7275 - accuracy:
0.7445 - val_loss: 0.9774 - val_accuracy: 0.6732
Epoch 7/20
50000/50000 [=====] - 64s 1ms/sample - loss: 0.6469 - accuracy:
0.7720 - val_loss: 0.9624 - val_accuracy: 0.6800
Epoch 8/20
50000/50000 [=====] - 67s 1ms/sample - loss: 0.5795 - accuracy:
0.7960 - val_loss: 1.0283 - val_accuracy: 0.6739
Epoch 9/20
50000/50000 [=====] - 65s 1ms/sample - loss: 0.5209 - accuracy:
0.8150 - val_loss: 1.0656 - val_accuracy: 0.6741

```

```
Out[33]: <tensorflow.python.keras.callbacks.History at 0x24d4358c888>
```

Check model performance

```
In [37]: metrics = pd.DataFrame(model.history.history)
```

```
In [38]: metrics.columns
```

```
Out[38]: Index(['loss', 'accuracy', 'val_loss', 'val_accuracy'], dtype='object')
```


```
In [39]: metrics.head()
```

```
Out[39]:
```

	loss	accuracy	val_loss	val_accuracy
0	1.506543	0.45526	1.268239	0.5467
1	1.165329	0.58668	1.093938	0.6142
2	1.006639	0.64738	1.083885	0.6196
3	0.898590	0.68648	1.061140	0.6235
4	0.801830	0.71856	0.995604	0.6623


```
In [40]: metrics[['loss', 'val_loss']].plot()
```

```
Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x24d001a2648>
```



```
In [43]: metrics[['accuracy', 'val_accuracy']].plot()
```

```
Out[43]: <matplotlib.axes._subplots.AxesSubplot at 0x24d00222208>
```



```
In [44]: #We stopped the training based on the Loss and not on the accuracy
```

Compare the same result using evaluate

```
In [45]: model.evaluate(X_test,y_cat_test,verbose=0)
```

```
Out[45]: [1.0655978028297424, 0.6741]
```

Do model prediction

```
In [50]: predictions = model.predict_classes(X_test)
```

Evaluation metrics and Reporting

```
In [51]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [53]: print(classification_report(y_test,predictions))
print('\n')
print(confusion_matrix(y_test,predictions))
```

precision recall f1-score support


0	0.70	0.73	0.71	1000
1	0.80	0.79	0.80	1000
2	0.52	0.63	0.57	1000
3	0.48	0.50	0.49	1000
4	0.63	0.61	0.62	1000
5	0.53	0.60	0.56	1000
6	0.77	0.74	0.75	1000
7	0.77	0.69	0.73	1000
8	0.81	0.77	0.79	1000
9	0.84	0.68	0.75	1000
accuracy			0.67	10000
macro avg	0.69	0.67	0.68	10000
weighted avg	0.69	0.67	0.68	10000

```
[[726 25 83 21 31 8 10 15 65 16]
 [ 32 793 25 9 6 11 15 10 31 68]
 [ 58 5 630 72 73 70 53 27 10 2]
 [ 17 12 101 500 66 197 58 34 9 6]
 [ 25 3 120 84 611 56 38 54 7 2]
 [ 13 6 82 186 39 599 17 43 8 7]
 [ 3 3 72 85 36 50 738 7 3 3]
 [ 20 1 42 39 78 107 7 694 5 7]
 [ 93 36 35 20 10 8 7 4 768 19]
 [ 53 105 26 23 17 25 13 17 39 682]]
```

Visualize confusion matrix

```
In [55]: plt.figure(figsize=(12,6))
sns.heatmap(confusion_matrix(y_test,predictions),annot=True)
```

```
Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x24d01d76708>
```




Predicting single image

```
In [56]: new_img = X_test[95]
```



```
plt.imshow(new_img)
```

Out[56]: <matplotlib.image.AxesImage at 0x24d43415288>



```
In [58]: y_test[95]
```

Out[58]: array([6], dtype=uint8)


```
In [60]: model.predict_classes(new_img.reshape(1,32,32,3))
```

Out[60]: array([6], dtype=int64)

So from above, it shows that our model predicted the frog correctly

```
In [61]: new_img2 = X_test[0]
plt.imshow(new_img2)
```

Out[61]: <matplotlib.image.AxesImage at 0x24d026dbc88>



```
In [62]: y_test[0]
```


Out[62]: array([3], dtype=uint8)

```
In [63]: model.predict_classes(new_img2.reshape(1,32,32,3))
```

Out[63]: array([3], dtype=int64)

```
In [65]: new_img3 = X_test[16]  
plt.imshow(new_img3)
```

```
Out[65]: <matplotlib.image.AxesImage at 0x24d024170c8>
```



```
In [66]: y_test[16]
```

```
Out[66]: array([5], dtype=uint8)
```

```
In [67]: model.predict_classes(new_img3.reshape(1,32,32,3))
```

```
Out[67]: array([5], dtype=int64)
```

```
In [ ]:
```